

# ElasticBF: Fine-grained and Elastic Bloom Filter Towards Efficient Read for LSM-tree-based KV Stores

Yueming Zhang, Yongkun Li, Fan Guo, Cheng Li, Yinlong Xu

Contact: [ykli@ustc.edu.cn](mailto:ykli@ustc.edu.cn)

*University of Science and Technology of China*

# Outline

## ➤ Background

- ❑ Key-value (KV) stores and LSM tree

## ➤ Motivation

- ❑ Read amplification problem in KV stores

## ➤ Design of ElasticBF

## ➤ Performance Evaluation

## ➤ Conclusion

# Background

➤ Key-value (KV) store has become an important storage engine for many applications

- ❑ Cloud systems
- ❑ Social networks
- ❑ ...

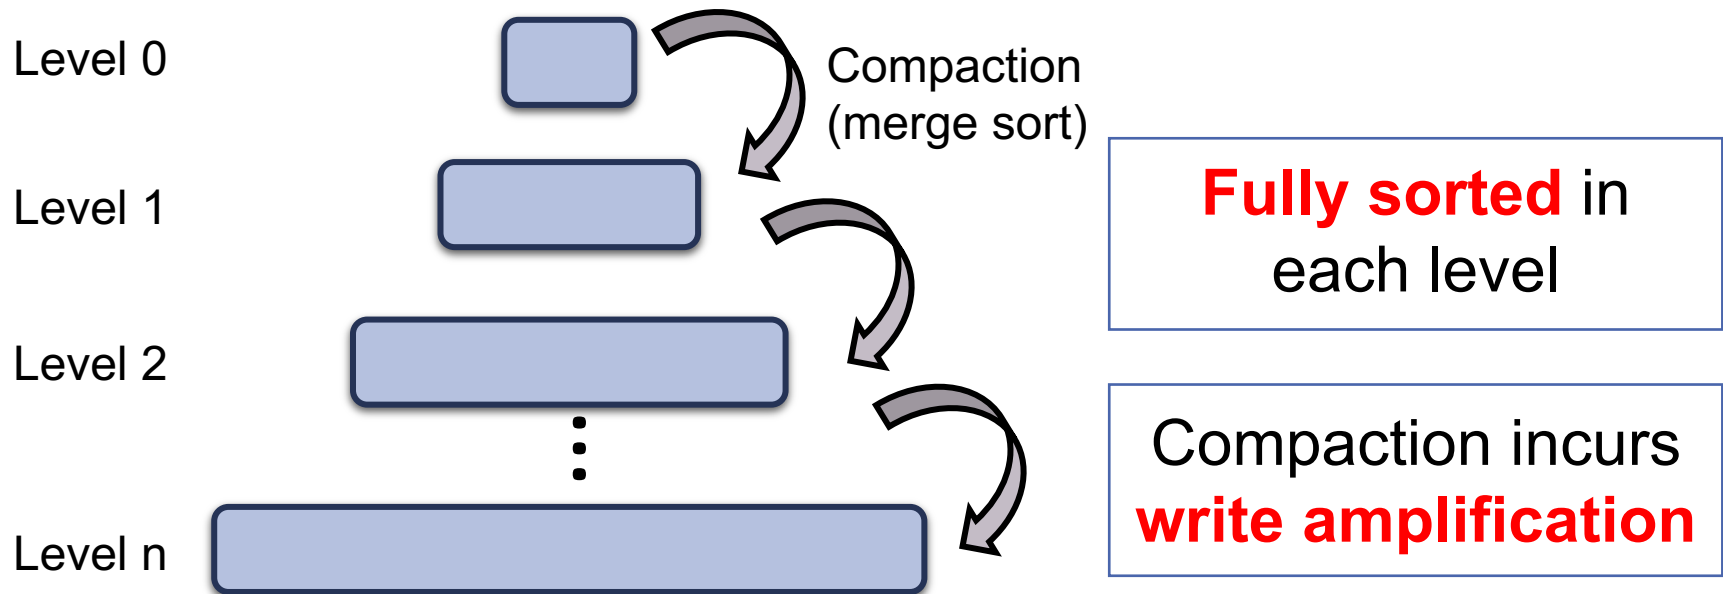
➤ Examples of KV stores

- ❑ Hbase @ Apache
- ❑ LevelDB @ Google
- ❑ RocksDB @ Facebook
- ❑ ...



# LSM Tree


- The most common design of KV stores is based on LSM-tree (log structured merge tree)



Data is written to Level 0 first, then merged to Level 1 via compaction, then Level 2, and so on.

# LSM Tree

- The most common design of KV stores is based on LSM-tree (log structured merge tree)




Looking up a key requires multiple I/O requests as it may require to search in multiple levels (**read amplification**).

# LevelDB

## ➤ One typical implementation of LSM tree

- ❑ Focus on data layout on disk




It suffers from **read amplification problem**, especially for a large KV store which has multiple levels.

# Bloom Filter

## ➤ Bloom filter in each SSTable

- ❑ A bit array with multiple hash functions
- ❑ Help quickly identify whether a key exists in an SSTable or not



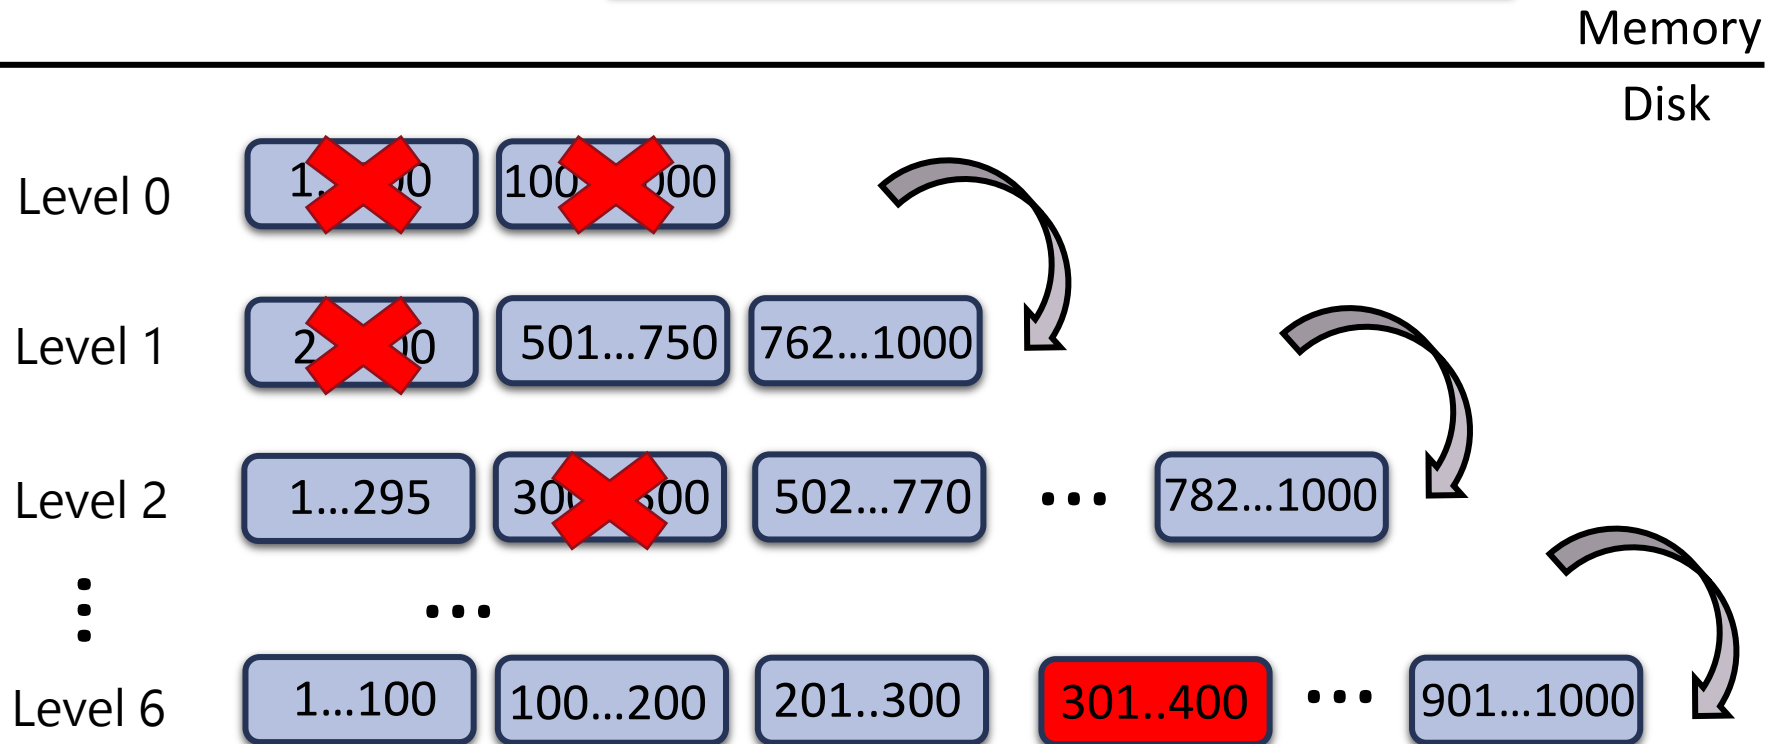
- ❑ Bloom filter suffers from **false positive** (hash collision)
  - False positive rate (FPR) :  $0.6185^b$  (b: Bits-per-key)

Bits-per-key	2bits	3bits	4bits	5bits	6bits
FPR	40%	23.7%	14.7%	9.2%	5.6%

# Read Flow with Bloom Filter

Example: Get(301)

Cached Bloom filters



**Bloom filters are required to be cached in memory**




# Motivation

- KV stores suffer from large read amplification
  - ❑ Bloom filter reduces read I/O, but has false positive
  - ❑ Reducing false positive may need to allocate many bits for each key, incurs large memory overhead
  
- Question: how to improve the Bloom filter design with limited memory consumption so as to
  - ❑ reduce extra I/O requests and
  - ❑ improve read performance of KV stores

# Main Idea

## ➤ Observation

- ❑ Access frequencies of SSTables in low levels are higher
- ❑ Unevenness of accesses even within the same level




## ➤ Main idea: ElasticBF

- ❑ An elastic scheme according to access frequency
- ❑ SSTables with high (low) access frequency
  - More (less) powerful Bloom filter (i.e., more (fewer) bits per key)
  - Lower (higher) false positive rate: fewer extra I/Os
  - Larger (smaller) memory consumption

# ElasticBF Design

## ➤ Challenge to realize an elastic scheme according to access frequency of SSTables

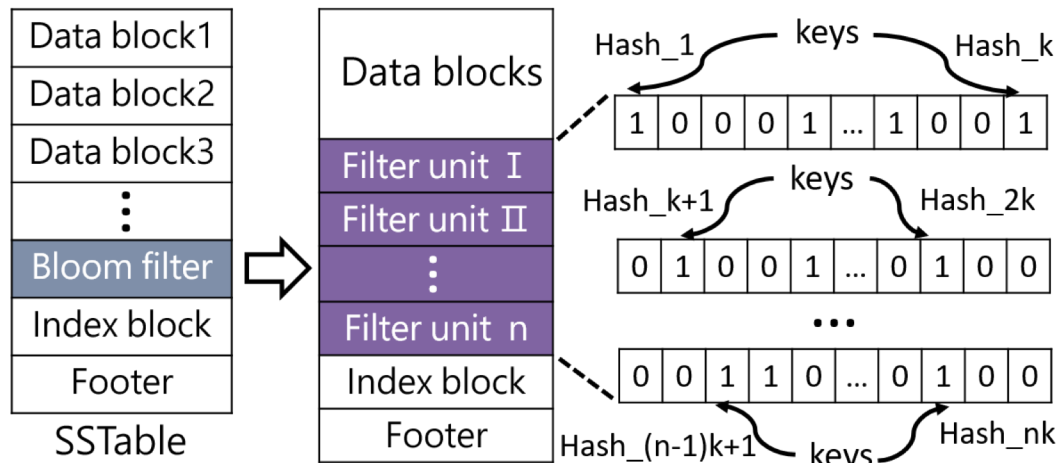
- ❑ Data organization in SSTable is fixed after creation
- ❑ Adjusting the Bloom filter in SSTables requires to reorganize the data



# ElasticBF Design

## ➤ Choice of ElasticBF: Step 1

- ❑ Build multiple small filter units in each SSTable with different and independent hash functions




## ➤ Rationale

- ❑ **Separability**: Multiple filters have the same FPR as a single filter with the same *bits-per-key*
  - FRR of n filter units :  $\prod_{i=1}^n 0.6185^{b_i} = 0.6185^b \quad (\sum_{i=1}^n b_i = b)$

# ElasticBF Design

## ➤ Choice of ElasticBF: Step 2

- ❑ Dynamically adjust the filter units in memory for each SSTable according to its access frequency
  - Enable more filter units by loading them into memory
  - Disable in-memory filter units by simply discarding them



Elastic feature: false positive rate can be dynamically adjusted

Data organization in SSTable does not change

# Key Issues

**How to determine the most appropriate number of filter units for each SSTable ?**

Adjusting Rule

**How to realize a dynamic adjustment with small overhead ?**

Multi-Queue

# Adjusting Rule

## ➤ Goal

- ❑ Try to reduce the extra I/Os caused by false positive


$$E[Extra\_IO] = \sum_{i=1}^N fp_i * f_i$$

- ❑ Access frequency of SSTable  $i$  :  $f_i$
- ❑ False positive rate of the Bloom filter in SSTable  $i$  :  $fp_i$
- ❑ Number of SSTables in the KV store :  $N$

ElasticBF estimates  $f_i$  in the runtime and adjusts  $fp_i$  accordingly so as to minimize  $E[Extra\_IO]$

# Multi-Queue


- Guides dynamic adjustment of the number of enabled filter units for each SSTable
  - ❑ Multiple least recently used queues ( LRU )
  - ❑  $Q_i$  corresponds to the SSTables with  $i$  filter units being enabled ( $Q_n$  : hottest SSTables,  $Q_0$  : coldest SSTables)





# Multi-Queue

- Dynamically adjust the filter units in Multi-Queue
  - ❑ Enable filter unit when the SSTable is accessed and  $E[Extra\_IO]$  can be reduced
  - ❑ Disable filter unit according to expiring policy



# Overhead Analysis

## ➤ Storage overhead

Size of KV pair	Size of SSTable	# KV pairs in a SSTable	bits-per-key	Space percent
1KB	2MB	2048	4	0.05%

## ➤ Computation overhead

- ❑ Time of building filters : ~1%
- ❑ Sufficient CPU resources
  - Multi-threading: generate multiple filter units simultaneously

## ➤ Memory overhead

Size of database	Number of SSTables	Memory overhead
100GB	50K	200KB

# Experiment Setting

## ➤ Experiment environment

### ❑ Machine

CPU	Disk	OS
Intel(R) Xeon(R) E5-2650 v4 @ 2.20GHz	Intel 3700 series SSD	CentOS 7.0/ Linux 3.10.0-5.14

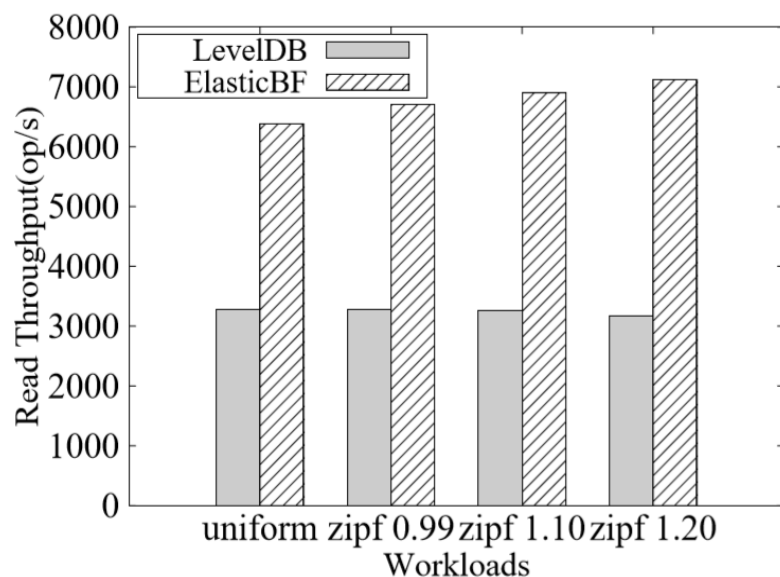
### ❑ Workloads: YCSB

Size of KV pair	Size of database	Request Distribution	Zipfian skew	Zero lookup/ Non-zero lookup	Number of Get Requests
1024	100 GB	zipfian/uniform	0.99/1.1/1.2	1:1	1 million

# Experiment Results

## ➤ Read Performance

Throughput




Number of I/Os for data access

	uniform	zipf 0.99	zipf 1.10	zipf 1.20
LevelDB	1525595	1585605	1634752	1667947
ElasticBF	628225	578553	550658	545345

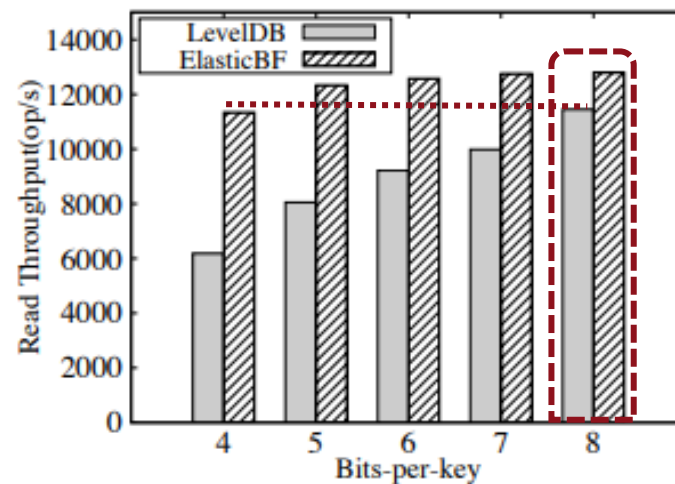
ElasticBF can achieve  $1.94\times$ - $2.24\times$  read throughput and greatly reduce the number of I/Os for data access compared to LevelDB

# Experiment Results

## ➤ Read Performance vs Memory Usage



(a) Uniform Workload



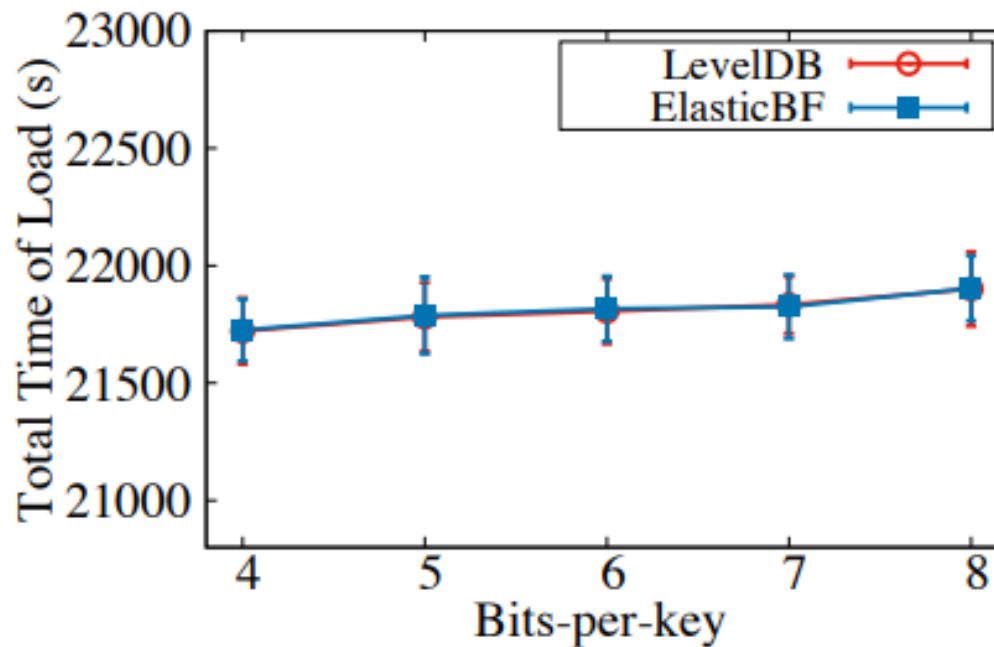
(b) Zipf0.99 Workload

ElasticBF can achieve a similar read performance with LevelDB with only a half memory usage

# Experiment Results

## ➤ Write Performance

❑ Load 100GB KV store



ElasticBF has almost the same write throughput with LevelDB

# Conclusion

- LSM tree suffers from read amplification problem
  - ❑ Bloom filter reduces extra I/Os during read
  - ❑ Uniform Bloom filter design either suffers from high false positive rate or incurs large memory overhead
- We develop ElasticBF
  - ❑ An elastic scheme to dynamically adjust the Bloom filters in SSTables according to access frequency
  - ❑ Improves read performance with limited memory
  - ❑ Orthogonal to works optimizing LSM-tree structure

# Thanks for your attention!

*For any questions, please feel free to contact  
Prof. Yongkun Li at USTC.*

[ykli@ustc.edu.cn](mailto:ykli@ustc.edu.cn)

<http://staff.ustc.edu.cn/~ykli/>